



 Strategic — platform default (most carry a condition)
 Tactical — a stated niche or estate
 Experimental — pilot only, outside regulated paths
 AWS Strategic only where that cloud is primary

C1 AI traffic gateway model, tool (MCP) and agent (A2A) calls

The diagram illustrates AI security solutions categorized by three main areas: secrets, supply chain, and sandboxing. These categories are represented by blue rounded rectangles at the top. Below each category, specific solutions are listed in white rounded rectangles with blue borders. The solutions are arranged in a grid-like fashion, with some spanning multiple categories.

Category	Solution	Description
AI security	HashiCorp Vault	IBM · agentic IAM · BUSL
AI security	Model & package scanning	safetensors by default
AI security	OpenSSF Model Signing	sign weights you produce
secrets	Prisma AIRS	Palo Alto estates
secrets	HiddenLayer	independent specialist
secrets	Check Point AI Guardrails	runtime detector · self-host for client data

OpenLineage + firm GenAI facets	IBM watsonx.governance IBM estates	Credo AI policy-led programmes
Collibra AI Governance where Collibra is the catalogue	ModelOp after full due diligence	ValidMind model-risk-led firms

PATTERN - NOT SCORED Primary cloud's model service Bedrock · Foundry · Gemini Enterprise Agent Platform, in region	vLLM default engine · the exit route	Fireworks AI once ISO certificates confirmed	SQLang backup engine once CVE is fixed
Hugging Face Hub governed open-weight supply	Together AI EU dedicated, ZDR on	OpenRouter behind the gateway · Stripe deal	Ollama developer tier
Cerebras low latency, non-confidential	LM Studio desktops only · no service use	llm-d CNCF sandbox · pilot	NVIDIA Dynamo beta · pilot

The diagram illustrates a collection of AI agents and their specific functions, organized into a grid-like structure. The agents are represented by colored boxes with rounded corners, each containing a title and a brief description of its capabilities. The agents are: A2A 1.0 (cross-team delegation, signed cards), MCP (behind a governed gateway, alt: OpenAPI), AWS AgentCore Gateway + Identity (managed tool governance), Browserbase (only where no API exists), E2B (microVM sandbox), Agent Skills (script-free, internal, alt: AGENTS.md), Exa (via egress proxy + DLP), Tavily (Nebius-owned), and Composio (May 2026 token incident, self-host). The Composio box is highlighted with a dashed orange border.

Agent	Capabilities
A2A 1.0	cross-team delegation · signed cards
MCP	behind a governed gateway · alt: OpenAPI
AWS AgentCore Gateway + Identity	managed tool governance
Browserbase	only where no API exists
E2B	microVM sandbox
Agent Skills	script-free, internal · alt: AGENTS.md
Exa	via egress proxy + DLP
Tavily	Nebius-owned
Composio	May 2026 token incident · self-host

Dooling default engine · LF AI & Data	Unstructured ACL-aware connectors	Google Cloud Google Document AI processor region confirmed	Mistral OCR 4.1 pin the model ID
LlamaParse parse and extract only	Firecrawl AGPL server · cloud with ZDR	Reducto hard documents	Apify public data only
Crawl4AI pre-1.0 · pilots	MinerU licence thresholds apply		

- Tiers are for a regulated asset manager, scored on one eight-criterion rubric with regulated-FS weights. A Strategic tier carries its condition; the condition is the decision.
- Cloud-tagged items are alternatives chosen by primary cloud, not a shopping list.
- Disclosure: researched and drafted with an Anthropic model. Anthropic items were scored on the same rubric; their tiers were set by the reader, and an independent alternative is named for each.
- Personal research, not any firm's platform. Every tile is backed by sourced facts in the dataset and the product appendix.