

You may not need a vector database

Step 0 Non-negotiables, whatever the store

- entitlement filters applied inside the search
- no unfiltered fallback
- rebuild from source tested twice a year



Step 1 Use what you run

relational database search engine document database
object-store tier



Step 2 Trigger shown in a load test on your own data?

- (a) filtered latency or recall fails the budget
- (b) tenants must be isolated at a scale schemas cannot manage
- (c) corpus makes the existing platform's cost disproportionate

no trigger stay in Step 1



at least one: yes

Step 3 Dedicated engine

chosen by where it must run

LEAKAGE · ILLUSTRATIVE

Filter after ranking + fallback

rank the top 20 passages



filter out other funds: none left



retry without the filter



another client's material reaches the draft

SAFE

Filter inside search + hard empty result

only permitted passages are candidates



none relevant: return nothing



draft flags "no comparable commentary"

THE SIGNAL

Filtered recall@k

recall measured with production-like entitlement filters applied, not on the open index